

Exercise - Taints and Tolerations

Name: Haffiz Hissham

Date: 21 September 2024

Q1

create a taint on node01 that not allow any new pod schedule on node.

node01 is the name of the node you want to taint.

key=value is the key-value pair for the taint.

in this case use key=node-restriction and value=true.

Run a pod pod1 with image=nginx that use the toleration.

Make sure pod is running .

```
GNU nano 4.8
apiVersion: v1
kind: Pod
metadata:
  name: pod1
spec:
  containers:
  - name: nginx
    image: nginx
    imagePullPolicy: IfNotPresent
  tolerations:
  - key: "node-restriction"
    operator: "Equal"
    value: "true"
    effect: "NoSchedule"
```

```
controlplane $
controlplane $ kubectl taint nodes node01 node-restriction=true:NoSchedule
node/node01 tainted
controlplane $
controlplane $ kubectl get nodes -o='custom-columns=NodeName:.metadata.name,TaintKey:.spec.taints[*].key,TaintValue:.spec.taints[*].value,TaintEffect:.spec.taints[*].effect'
NodeName      TaintKey      TaintValue    TaintEffect
controlplane  node-role.kubernetes.io/control-plane  <none>      NoSchedule
node01        node-restriction  true         NoSchedule
controlplane $
controlplane $ nano pod1.yaml
controlplane $
controlplane $ k apply -f pod1.yaml
pod/pod1 created
controlplane $
controlplane $ k get pods -o wide
NAME    READY   STATUS    RESTARTS   AGE   IP        NODE    NOMINATED NODE   READINESS GATES
pod1    0/1     ContainerCreating   0       7s    <none>    node01   <none>           <none>
controlplane $
controlplane $ k get pods -o wide
NAME    READY   STATUS    RESTARTS   AGE   IP        NODE    NOMINATED NODE   READINESS GATES
pod1    1/1     Running   0          11s   192.168.1.4   node01   <none>           <none>
controlplane $
controlplane $
```

```
Node-Selectors:  <none>
Tolerations:    node-restriction=true:NoSchedule
                node.kubernetes.io/not-ready:NoExecute op=Exists for 300s
                node.kubernetes.io/unreachable:NoExecute op=Exists for 300s

Events:
  Type    Reason      Age   From          Message
  ----    -
  Normal  Scheduled   54s   default-scheduler  Successfully assigned default/pod1 to node01
  Normal  Pulling     54s   kubelet         Pulling image "nginx"
  Normal  Pulled      44s   kubelet         Successfully pulled image "nginx" in 9.585s (9.585s including waiting). Image size: 71027698 bytes.
  Normal  Created     44s   kubelet         Created container nginx
  Normal  Started     44s   kubelet         Started container nginx
controlplane $
```

Q2

Create a single Pod of image httpd:2.4.41-alpine in Namespace default.

The Pod should be named pod1 and the container should be named pod1-container. This Pod should only be scheduled on controlplane nodes.

Do not add new labels to any nodes.

```
apiVersion: v1
kind: Pod
metadata:
  creationTimestamp: null
  labels:
    run: pod1
  name: pod1
spec:
  nodeName: controlplane
  containers:
  - image: httpd:2.4.41-alpine
    name: pod1-container
    resources: {}
  dnsPolicy: ClusterFirst
  restartPolicy: Always
status: {}
```

```
controlplane $
controlplane $ k run pod1 --image=httpd:2.4.41-alpine --dry-run=client -o yaml > q2.yaml
controlplane $ nano q2.yaml
controlplane $
controlplane $ k apply -f q2.yaml
pod/pod1 created
controlplane $
controlplane $ k get pods -o wide

```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED NODE	READINESS GATES
pod1	1/1	Running	0	11s	192.168.0.4	controlplane	<none>	<none>

```
controlplane $
controlplane $
```