

MOCK EXAM 1.0

Name: Haffiz Hissham

Date: 22 September 2024

Q1

Create a nginx pod called dns-resolver using image nginx expose it internally with a service called dns-resolver-service.

check if pod and service name are resolvable from within the cluster.

use the image: busybox:1.28 for dns lookup

save the result in /root/nginx.svc.

```
controlplane $ k run dns-resolver --image=nginx
pod/dns-resolver created
controlplane $ k get pods
NAME          READY   STATUS    RESTARTS   AGE
dns-resolver  1/1     Running   0           54s
controlplane $ kubectl expose pods dns-resolver --port=80 --name=dns-resolver-service
service/dns-resolver-service exposed
controlplane $ k get svc
NAME          TYPE          CLUSTER-IP   EXTERNAL-IP   PORT(S)   AGE
dns-resolver-service  ClusterIP    10.105.221.87 <none>        80/TCP    4s
kubernetes    ClusterIP    10.96.0.1    <none>        443/TCP   12d
controlplane $ k delete svc dns-resolver-service
service "dns-resolver-service" deleted
controlplane $ kubectl expose pods dns-resolver --port=80 --name=dns-resolver-service --type=ClusterIP
service/dns-resolver-service exposed
controlplane $ k delete svc dns-resolver-service
service "dns-resolver-service" deleted
controlplane $ kubectl expose pods dns-resolver --port=80 --name=dns-resolver-service --type=ClusterIP
service/dns-resolver-service exposed
controlplane $ k get svc
NAME          TYPE          CLUSTER-IP   EXTERNAL-IP   PORT(S)   AGE
dns-resolver-service  ClusterIP    10.111.202.77 <none>        80/TCP    3s
kubernetes    ClusterIP    10.96.0.1    <none>        443/TCP   12d
controlplane $
```

```
controlplane $ k run ns-pod --image=busybox:1.28 --rm -it -- nslookup dns-resolver-service
pod "ns-pod" deleted
error: timed out waiting for the condition
controlplane $ k run ns-pod --image=busybox:1.28 --rm -it --restart=Never -- nslookup dns-resolver-service
Server:      10.96.0.10
Address 1: 10.96.0.10 kube-dns.kube-system.svc.cluster.local

Name:      dns-resolver-service
Address 1: 10.111.202.77 dns-resolver-service.default.svc.cluster.local
pod "ns-pod" deleted
controlplane $ k run ns-pod --image=busybox:1.28 --rm -it --restart=Never -- nslookup dns-resolver-service > /root/nginx.svc
controlplane $ cat /root/nginx.svc
Server:      10.96.0.10
Address 1: 10.96.0.10 kube-dns.kube-system.svc.cluster.local

Name:      dns-resolver-service
Address 1: 10.111.202.77 dns-resolver-service.default.svc.cluster.local
pod "ns-pod" deleted
controlplane $
```

Q2

Create a persistent volume with name app-data, of capacity 2Gi and access mode ReadOnlyMany.

The type of volume is hostPath and its location is /srv/app- data

```
GNU nano 4.8
apiVersion: v1
kind: PersistentVolume
metadata:
  name: app-data
  labels:
    type: local
spec:
  storageClassName: manual
  capacity:
    storage: 2Gi
  accessModes:
    - ReadOnlyMany
  hostPath:
    path: "/srv/app- data"
```

```
controlplane $
controlplane $ nano app-data.yaml
controlplane $
controlplane $ k apply -f app-data.yaml
persistentvolume/app-data created
controlplane $
controlplane $ k get pv
NAME          CAPACITY  ACCESS MODES  RECLAIM POLICY  STATUS   CLAIM   STORAGECLASS  VOLUMEATTRIBUTESCLASS  REASON  AGE
app-data      2Gi       ROX           Retain          Available             manual      <unset>              4s
controlplane $
```

Q3

Check to see how many nodes are ready (not including nodes tainted NoSchedule) and write the

number to /opt/KUSC00402/kusc00402.txt.

```
controlplane $
controlplane $ JSONPATH='{range .items[*]}{@.metadata.name}:{range @.status.conditions[*]}{@.type}={@.status};{end}-{end}\' \
> && kubectl get nodes -o jsonpath="$JSONPATH" | grep "Ready=True"
controlplane:NetworkUnavailable=False;MemoryPressure=False;DiskPressure=False;PIDPressure=False;Ready=True;node01:NetworkUnavaila
ble=False;MemoryPressure=False;DiskPressure=False;PIDPressure=False;Ready=True;
controlplane $
controlplane $ mkdir -p /opt/KUSC00402
controlplane $
controlplane $ echo "1" > /opt/KUSC00402/kusc00402.txt
controlplane $
controlplane $ cat /opt/KUSC00402/kusc00402.txt
1
controlplane $
```

Q4

Create a new pod called mock-pod with image busy box
Allow the pod to be able to set system_time

The container should sleep for 4000 seconds

```
GNU nano 4.8
apiVersion: v1
kind: Pod
metadata:
  name: mock-pod
spec:
  containers:
  - name: mock-pod
    image: busybox
    command: [ "sh", "-c", "sleep 4000" ]
    securityContext:
      capabilities:
        add: ["SYS_TIME"]

controlplane $
controlplane $ nano mock-pod.yaml
controlplane $
controlplane $ k apply -f mock-pod.yaml
pod/mock-pod created
controlplane $
controlplane $ k get pods
NAME           READY   STATUS    RESTARTS   AGE
dns-resolver   1/1     Running   0           13m
mock-pod       1/1     Running   0           12s
controlplane $
```

```
IP: 192.168.1.8
IPs:
  IP: 192.168.1.8
Containers:
  mock-pod:
    Container ID: containerd://4e89a1c54a3c7b2be60a54
    Image: busybox
    Image ID: docker.io/library/busybox@sha256:c2
    Port: <none>
    Host Port: <none>
    Command:
      sh
      -c
      sleep 4000
    State: Running
      Started: Sun, 22 Sep 2024 02:25:04 +0000
    Ready: True
    Restart Count: 0
    Environment: <none>
    Mounts:
```

Q5

Temporarily stop the kube-scheduler, this means in a way that you can start it again afterwards.

Create a single Pod named manual-schedule of image httpd:2.4-alpine, confirm it's created but not scheduled on any node.

Now you're the scheduler and have all its power, manually schedule that Pod on node with nodename. Make sure it's running.

Start the kube-scheduler again and confirm it's running correctly by creating a second Pod named manual-schedule2 of image httpd:2.4-alpine on controlplane

```
GNU nano 4.8 /etc/kubernetes/manifests/kube-scheduler.yaml
apiVersion: v1
kind: Pod
metadata:
  creationTimestamp: null
  labels:
    component: kube-scheduler
    tier: control-plane
  name: kube-scheduler
  namespace: kube-system
spec:
  containers:
  - command:
    - kube-scheduler
    - --authentication-kubeconfig=/etc/kubernetes/scheduler.conf
    - --authorization-kubeconfig=/etc/kubernetes/scheduler.conf
    - --bind-address=127.0.0.1
    - --kubeconfig=/etc/kubernetes/scheduler.conf
    - --leader-elect=true
    image: registry.k8s.io/kube-scheduler:v1.30.0
    name: kube-scheduler
    resources: {}
  dnsPolicy: ClusterFirst
  restartPolicy: Always
  schedulerName: default-scheduler
```

```
controlplane $ k get pods -A | grep scheduler
kube-system      kube-scheduler-controlplane      1/1      Running    2 (41m ago)    12d
controlplane $
controlplane $
controlplane $
controlplane $ k get pods -A -o wide | grep scheduler
kube-system      kube-scheduler-controlplane      1/1      Running    2 (41m ago)    12d    172.30.1.2    controlplane
e <none>         <none>
controlplane $ nano /etc/kubernetes/manifests/kube-scheduler.yaml
controlplane $
controlplane $ k get pods -A -o wide | grep scheduler
kube-system      kube-scheduler-controlplane      1/1      Running    2 (42m ago)    12d    172.30.1.2    controlplane
e <none>         <none>
controlplane $
controlplane $ k run manual-schedule --image=httpd:2.4-alpine
pod/manual-schedule created
controlplane $
controlplane $ k get pods -o wide | grep manual
manual-schedule  0/1      Pending    0          14s    <none>      <none>      <none>      <none>
controlplane $
```

```
GNU nano 4.8
apiVersion: v1
kind: Pod
metadata:
  creationTimestamp: "2024-09-22T03:46:02Z"
  labels:
    run: manual-schedule
  name: manual-schedule
  namespace: default
  resourceVersion: "5759"
  uid: 88d1bf73-09aa-4c44-91dd-21461bd146ee
spec:
  nodeName: node01
  containers:
  - image: httpd:2.4-alpine
    imagePullPolicy: IfNotPresent
    name: manual-schedule
    resources: {}
    terminationMessagePath: /dev/termination-log
```

```
controlplane $
controlplane $ k get pods -o wide | grep manual
manual-schedule 0/1 Pending 0 87s <none> <none> <none> <none>
controlplane $ k delete pods manual-schedule
pod "manual-schedule" deleted
controlplane $
controlplane $ k get pods -o wide | grep manual
controlplane $
controlplane $ nano manual.yaml
controlplane $
controlplane $
controlplane $
controlplane $
controlplane $
controlplane $ k apply -f manual.yaml
pod/manual-schedule created
controlplane $
controlplane $ k get pods -o wide | grep manual
manual-schedule 1/1 Running 0 7s 192.168.1.7 node01 <none> <none>
controlplane $
```

```
Exam Desktop Editor 1001
GNU nano 4.8
apiVersion: v1
kind: Pod
metadata:
  creationTimestamp: "2024-09-22T03:46:02Z"
  labels:
    run: manual-schedule2
  name: manual-schedule2
  namespace: default
  resourceVersion: "5759"
  uid: 88d1bf73-09aa-4c44-91dd-21461bd146ee
spec:
  nodeName: controlplane
  containers:
  - image: httpd:2.4-alpine
    imagePullPolicy: IfNotPresent
    name: manual-schedule2
    resources: {}
    terminationMessagePath: /dev/termination-log
    terminationMessagePolicy: File
    volumeMounts:
    - mountPath: /var/run/secrets/kubernetes.io/serviceaccount
      name: kube-api-access-djjff
```

```

controlplane $
controlplane $
controlplane $ nano /etc/kubernetes/manifests/kube-scheduler.yaml
controlplane $ cp manual.yaml manual2.yaml
controlplane $
controlplane $ nano manual2.yaml
controlplane $
controlplane $ k apply -f manual2.yaml
pod/manual-schedule2 created
controlplane $
controlplane $ k get pods -o wide | grep manual
manual-schedule      1/1      Running      0          2m21s    192.168.1.7    node01    <n
manual-schedule2    0/1      ContainerCr  0          6s       <none>        controlplane <n
controlplane $

controlplane $
controlplane $ k get pods -o wide | grep manual
manual-schedule      1/1      Running      0          3m8s     192.168.1.7    node01    <none>    <none>
manual-schedule2    1/1      Running      0          53s      192.168.0.5    controlplane <none>    <none>
controlplane $
controlplane $
controlplane $

```

Q6

Create a pod called pod-cka with two containers, as given below:

Container 1 - name: cka2, image: nginx

Container2 - name: cka2, image:

busybox,

command: sleep 3000

```

GNU nano 4.8
apiVersion: v1
kind: Pod
metadata:
  creationTimestamp: null
  labels:
    run: pod-cka
  name: pod-cka
spec:
  containers:
  - image: nginx
    name: cka1
  - image: busybox
    name: cka2
    command: [ "sh", "-c", "sleep 3000" ]
  dnsPolicy: ClusterFirst
  restartPolicy: Always
status: {}

controlplane $
controlplane $ k run pod-cka --image=cka1 --dry-run=client -o yaml > q6.yaml
controlplane $
controlplane $ nano q6.yaml
controlplane $
controlplane $ k apply -f q6.yaml
pod/pod-cka created
controlplane $
controlplane $ k get pods
NAME           READY   STATUS    RESTARTS   AGE
dns-resolver   1/1     Running   0           19m
mock-pod       1/1     Running   0           5m54s
pod-cka        2/2     Running   0           8s
controlplane $

```

Q7

create a deployment named source-ip-app that uses the image registry.k8s.io/echoserver:1.4 .

```
controlplane $  
controlplane $ k create deployment source-ip-app --image=registry.k8s.io/echoserver:1.4  
deployment.apps/source-ip-app created  
controlplane $  
controlplane $ k get deployments.apps  
NAME          READY   UP-TO-DATE   AVAILABLE   AGE  
source-ip-app  0/1     1            0           5s  
controlplane $  
controlplane $ k get deployments.apps  
NAME          READY   UP-TO-DATE   AVAILABLE   AGE  
source-ip-app  1/1     1            1           9s  
controlplane $
```

Q8

create a pod as follows:

name:mongo

using image:mongo

in a new kubernetes namespace named:my-website

```
controlplane $  
controlplane $ k create ns my-website  
namespace/my-website created  
controlplane $  
controlplane $ k get ns  
NAME          STATUS   AGE  
default       Active  12d  
kube-node-lease  Active  12d  
kube-public    Active  12d  
kube-system    Active  12d  
local-path-storage  Active  12d  
my-website     Active  3s  
controlplane $  
controlplane $ k run mongo --image=mongo -n my-website  
pod/mongo created  
controlplane $  
controlplane $ k get pods -n my-website  
NAME    READY   STATUS             RESTARTS   AGE  
mongo   0/1     ContainerCreating  0          10s  
controlplane $  
controlplane $ k get pods -n my-website  
NAME    READY   STATUS             RESTARTS   AGE  
mongo   0/1     ContainerCreating  0          13s  
controlplane $ k get pods -n my-website  
NAME    READY   STATUS    RESTARTS   AGE  
mongo   1/1     Running   0          17s  
controlplane $
```

Q9

You're asked to find out the following information about the cluster:

How many controlplane nodes are available?

How many worker nodes are available?

How many static pods are running?

Write your answers into file /opt/course/14/cluster-info, structured like this:

/opt/course/14/cluster-info

1: [ANSWER]

2: [ANSWER]

3: [ANSWER]

```
controlplane $
controlplane $ k get nodes
NAME                STATUS    ROLES    AGE   VERSION
controlplane        Ready    control-plane   12d   v1.30.0
node01              Ready    <none>         12d   v1.30.0
controlplane $
controlplane $ ls /etc/kubernetes/manifests/
etcd.yaml kube-apiserver.yaml kube-controller-manager.yaml kube-scheduler.yaml
controlplane $
controlplane $ ssh node01
Last login: Sun Nov 13 17:27:09 2022 from 10.48.0.33
node01 $
node01 $ ls etc/
ls: cannot access 'etc/': No such file or directory
node01 $ ls /etc/k
kernel/          kernel-img.conf  killercoda/      kubernetes/
node01 $ ls /etc/kubernetes/
kubelet.conf manifests/  pki/
node01 $ ls /etc/kubernetes/manifests/
node01 $ ls /etc/kubernetes/manifests/ -a
.  ..  .kubelet-keep
node01 $ ls /etc/kubernetes/manifests/ -l
total 0
controlplane $
controlplane $ nano /opt/course/14/cluster-info
controlplane $ cat /opt/course/14/cluster-info
# /opt/course/14/cluster-info
1: 1
2: 1
3: 4
controlplane $
```


Q10

Create a new deployment called mockpod, with image nginx:1.16 and 1 replica.

Next upgrade the deployment to version 1.17 using rolling update

Make sure that the version upgrade is recorded in the resource annotation

```
controlplane $  
controlplane $ k create deployment mockpod --image=nginx:1.16 --replicas=1  
deployment.apps/mockpod created  
controlplane $  
controlplane $ k get deployments.apps  
NAME           READY   UP-TO-DATE   AVAILABLE   AGE  
mockpod        1/1     1            1           37s  
source-ip-app  1/1     1            1           10m  
controlplane $  
controlplane $ kubectl set image deployment/mockpod nginx=nginx:1.17 --record  
Flag --record has been deprecated, --record will be removed in the future  
deployment.apps/mockpod image updated  
controlplane $  
controlplane $ kubectl rollout history deployment/mockpod  
deployment.apps/mockpod  
REVISION  CHANGE-CAUSE  
1          <none>  
2          kubectl set image deployment/mockpod nginx=nginx:1.17 --record=true  
  
controlplane $
```

Q11

write a command into /opt/course/100/cluster_events.sh which shows the latest events in the whole cluster, ordered by time (metadata.creationTimestamp). use kubectl for it.

now delete the kube-proxy pod running on node controlplane node and write the events this caused

into /opt/course/100/pod_kill.log.

```
controlplane $  
controlplane $ mkdir -p /opt/course/100  
controlplane $  
controlplane $ echo -e "kubectl get events --sort-by=.metadata.creationTimestamp" > /opt/course/100/cluster_events.sh  
controlplane $  
controlplane $ cat /opt/course/100/cluster_events.sh  
kubectl get events --sort-by=.metadata.creationTimestamp  
controlplane $  
controlplane $
```

```
controlplane $  
controlplane $ k get pods -A -o wide | grep controlplane  
kube-system      calico-kube-controllers-75bdb5b75d-zhhrq  1/1    Running  2 (77m ago)  12d    192.168.0.2    controlpla  
ne <none>         <none>                                     <none>                                     <none>  
kube-system      canal-szcfj                                2/2    Running  2 (77m ago)  12d    172.30.1.2    controlpla  
ne <none>         <none>                                     <none>                                     <none>  
kube-system      etcd-controlplane                         1/1    Running  2 (77m ago)  12d    172.30.1.2    controlpla  
ne <none>         <none>                                     <none>                                     <none>  
kube-system      kube-apiserver-controlplane               1/1    Running  2 (77m ago)  12d    172.30.1.2    controlpla  
ne <none>         <none>                                     <none>                                     <none>  
kube-system      kube-controller-manager-controlplane      1/1    Running  2 (77m ago)  12d    172.30.1.2    controlpla  
ne <none>         <none>                                     <none>                                     <none>  
kube-system      kube-proxy-mvqrk                         1/1    Running  2 (77m ago)  12d    172.30.1.2    controlpla  
ne <none>         <none>                                     <none>                                     <none>  
kube-system      kube-scheduler-controlplane               1/1    Running  2 (77m ago)  12d    172.30.1.2    controlpla  
ne <none>         <none>                                     <none>                                     <none>  
local-path-storage local-path-provisioner-75655fcf79-6xrsrw 1/1    Running  2 (77m ago)  12d    192.168.0.3    controlpla  
ne <none>         <none>                                     <none>                                     <none>  
controlplane $  
controlplane $  
controlplane $ k delete pods kube-proxy-mvqrk -n kube-system  
pod "kube-proxy-mvqrk" deleted  
controlplane $
```

```
controlplane $  
controlplane $ bash /opt/course/100/cluster_events.sh | tail -n 5 > /opt/course/100/pod_kill.log  
controlplane $  
controlplane $ cat /opt/course/100/pod_kill.log  
5m13s    Normal    Started      pod/mockpod-76cc984cd-7nc76      Started container nginx  
5m12s    Normal    SuccessfulDelete replicaset/mockpod-7686cdb85-kcqt Deleted pod: mockpod-7686cdb85-kcqt  
5m12s    Normal    ScalingReplicaSet deployment/mockpod                Scaled down replica set mockpod-7686cdb85 to 0 fr  
om 1  
5m12s    Normal    Killing      pod/mockpod-7686cdb85-kcqt       Stopping container nginx  
85s      Normal    Starting     node/controlplane  
controlplane $
```

Q12

create a deployment called pod-cka with two containers, as given below:

container 1 - name: cka1, image: nginx

container2 - name: cka2, image:busybox,

command: sleep 5000

```
GNU nano 4.8
apiVersion: v1
kind: Pod
metadata:
  creationTimestamp: null
  labels:
    run: pod-cka
  name: pod-cka
spec:
  containers:
  - image: nginx
    name: cka1
  - image: busybox
    name: cka2
    command: [ "sh", "-c", "sleep 5000" ]
  dnsPolicy: ClusterFirst
  restartPolicy: Always
status: {}

controlplane $
controlplane $ nano pod-cka.yaml
controlplane $
controlplane $ k apply -f pod-cka.yaml
pod/pod-cka created
controlplane $
controlplane $
controlplane $ k get pods | grep pod-cka
pod-cka                2/2      Running   0          10s
controlplane $
controlplane $
```

```
controlplane $ k describe pods pod-cka
Name:          pod-cka
Namespace:     default
Priority:       0
Service Account: default
Node:          node01/172.30.2.2
Start Time:    Sun, 22 Sep 2024 02:51:43 +0000
Labels:        run=pod-cka
Annotations:   cni.projectcalico.org/containerID: cc3269551ec410de5
               cni.projectcalico.org/podIP: 192.168.1.14/32
               cni.projectcalico.org/podIPs: 192.168.1.14/32
Status:        Running
IP:            192.168.1.14
IPs:
  IP: 192.168.1.14
Containers:
  cka1:
    Container ID:   containerd://6d91d8413036a65a3300f57a7f26d7ccd38ea
    Image:          nginx
    Image ID:       docker.io/library/nginx@sha256:04ba374043ccd2fc5c5
    Port:           <none>
    Host Port:      <none>
    State:          Running
      Started:      Sun, 22 Sep 2024 02:51:44 +0000
    Ready:          True
    Restart Count:  0
    Environment:    <none>
    Mounts:
      /var/run/secrets/kubernetes.io/serviceaccount from kube-api-acce
  cka2:
    Container ID:   containerd://627efb49276383317a2ef20f86e0aa0efe89d7
    Image:          busybox
    Image ID:       docker.io/library/busybox@sha256:c230832bd3b0be59a6
    Port:           <none>
    Host Port:      <none>
    Command:
      sh
      -c
      sleep 5000
    State:          Running
      Started:      Sun, 22 Sep 2024 02:51:45 +0000
    Ready:          True
```

Q13

use json path query to retrieve the osimages of all the nodes and store it in a file “all-nodes-os-info.txt” at root location.

note: the osimage are under the nodeInfo section under status of each node.

```
controlplane $  
controlplane $  
controlplane $ kubectl get node -o=jsonpath='{.items[*].status.nodeInfo.osImage}' > all-nodes-os-info.txt  
controlplane $  
controlplane $ cat all-nodes-os-info.txt  
Ubuntu 20.04.5 LTS Ubuntu 20.04.5 LTScontrolplane $
```

Q14

create a new persistentvolumeclaim:

- name: pv-volume
- class: csi-hostpath-sc
- capacity: 10mi

create a new pod which mounts the persistentvolumeclaim as a volume:

- name: web-server
- image: nginx
- mount path: /usr/share/nginx/html

configure the new pod to have readwriteonce access on the volume.

finally, using kubectl edit or kubectl patch expand the persistentvolumeclaim to a capacity of 70mi and record that change.

```
controlplane $  
controlplane $ nano q14-pv.yaml  
controlplane $  
controlplane $ k apply -f q14-pv.yaml  
persistentvolume/pv created  
persistentvolumeclaim/pv-volume created  
controlplane $  
controlplane $ k get pvc  
NAME          STATUS  VOLUME  CAPACITY  ACCESS MODES  STORAGECLASS  VOLUMEATTRIBUTESCLASS  AGE  
pv-volume     Bound   pv       10Mi      RWO           csi-hostpath-sc  <unset>                 3s  
controlplane $  
controlplane $
```

```

GNU nano 4.8
apiVersion: v1
kind: PersistentVolume
metadata:
  name: pv
  labels:
    type: local
spec:
  storageClassName: csi-hostpath-sc
  capacity:
    storage: 10Mi
  accessModes:
    - ReadWriteOnce
  hostPath:
    path: "/mnt/data"

---

apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: pv-volume
spec:
  storageClassName: csi-hostpath-sc
  accessModes:
    - ReadWriteOnce
  resources:
    requests:
      storage: 10Mi

```

```

GNU nano 4.8
apiVersion: v1
kind: Pod
metadata:
  name: web-server
spec:
  volumes:
    - name: task-pv-storage
      persistentVolumeClaim:
        claimName: pv-volume
  containers:
    - name: web-server
      image: nginx
      volumeMounts:
        - mountPath: "/usr/share/nginx/html"
          name: task-pv-storage

```

```

controlplane $
controlplane $ nano q14.yaml
controlplane $
controlplane $ k ap
api-resources (Print the supported API resources on the server)
api-versions (Print the supported API versions on the server, in the form of "group/version")
apply (Apply a configuration to a resource by file name or stdin)
controlplane $ k apply -f q14.yaml
pod/web-server created
controlplane $
controlplane $ k get pods
NAME          READY   STATUS    RESTARTS   AGE
web-server    0/1     ContainerCreating   0           4s
controlplane $
controlplane $ k describe pods web-server
Name:          web-server
Namespace:     default
Priority:       0
Service Account: default
Node:          node01/172.30.2.2
Start Time:    Sun, 22 Sep 2024 03:12:59 +0000
Labels:        <none>
Annotations:   cni.projectcalico.org/containerID: b2f2986540a74ced8fdea4863145f077aa24e75deca6e569
               cni.projectcalico.org/podIP: 192.168.1.4/32
               cni.projectcalico.org/podIPs: 192.168.1.4/32
Status:        Running
IP:            192.168.1.4
IPs:
  IP: 192.168.1.4
Containers:
  web-server:
    Container ID:  containerd://d205f49e07f4c62629ef23a02c43b921d4164b236f11f702862022f4c2230285
    Image:          nginx
    Image ID:       docker.io/library/nginx@sha256:04ba37404ccd2fc5c593885c0eacddeabd5ca375f9323666
    Port:          <none>
    Host Port:      <none>
    State:          Running
      Started:      Sun, 22 Sep 2024 03:13:10 +0000
    Ready:          True

```

```

apiVersion: v1
kind: PersistentVolume
metadata:
  annotations:
    kubect1.kubernetes.io/last-applied-configuration: |
      {"apiVersion":"v1","kind":"PersistentVolume","metadata":{"name":"pv","resourceVersion":"2842","uid":"dcf6675e-3834-4395-a984-b238b3add37c"},"spec":{"accessModes":["ReadWriteOnce"],"capacity":{"storage":"10Mi"},"claimRef":{"apiVersion":"v1","kind":"PersistentVolumeClaim","name":"pv-volume","namespace":"default","resourceVersion":"2840","uid":"1bec4f20-6fb3-4249-b84a-3523218e9572"},"storageClass":"csi-hostpath-sc","volumeMode":"Filesystem","volumeName":"pv"}}
  creationTimestamp: "2024-09-22T03:11:06Z"
  finalizers:
    - kubernetes.io/pv-protection
  labels:
    type: local
  name: pv
  resourceVersion: "2842"
  uid: dcf6675e-3834-4395-a984-b238b3add37c
spec:
  accessModes:
    - ReadWriteOnce
  capacity:
    storage: 10Mi
  claimRef:
    apiVersion: v1
    kind: PersistentVolumeClaim
    name: pv-volume
    namespace: default
    resourceVersion: "2840"
    uid: 1bec4f20-6fb3-4249-b84a-3523218e9572
  hostPath:
    path: /mnt/data
    type: ""

```

```

apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  annotations:
    kubect1.kubernetes.io/last-applied-configuration: |
      {"apiVersion":"v1","kind":"PersistentVolumeClaim","metadata":{"name":"pv-volume","namespace":"default","resourceVersion":"2844","uid":"1bec4f20-6fb3-4249-b84a-3523218e9572"},"spec":{"accessModes":["ReadWriteOnce"],"resources":{"requests":{"storage":"70Mi"}},"storageClass":"csi-hostpath-sc","volumeMode":"Filesystem","volumeName":"pv"}}
  creationTimestamp: "2024-09-22T03:11:06Z"
  finalizers:
    - kubernetes.io/pvc-protection
  name: pv-volume
  namespace: default
  resourceVersion: "2844"
  uid: 1bec4f20-6fb3-4249-b84a-3523218e9572
spec:
  accessModes:
    - ReadWriteOnce
  resources:
    requests:
      storage: 70Mi
  storageClassName: csi-hostpath-sc
  volumeMode: Filesystem
  volumeName: pv

```

```

controlplane $
controlplane $ kubectl patch -f q14-pv.yaml
error: must specify --patch or --patch-file containing the contents of the patch
controlplane $
controlplane $
controlplane $ k edit pv pv
persistentvolume/pv edited
controlplane $
controlplane $ k edit pvc pv-volume
error: persistentvolumeclaims "pv-volume" could not be patched: persistentvolumeclaims "pv-volume" is forbidden: only dynamically
provisioned pvc can be resized and the storageclass that provisions the pvc must support resize
You can run `kubectl replace -f /tmp/kubect1-edit-2547771406.yaml` to try this update again.
controlplane $ kubectl replace -f /tmp/kubect1-edit-2547771406.yaml
Error from server (Forbidden): error when replacing "/tmp/kubect1-edit-2547771406.yaml": persistentvolumeclaims "pv-volume" is fo
rbidden: only dynamically provisioned pvc can be resized and the storageclass that provisions the pvc must support resize
controlplane $

```

Q15

create a static pod named static-control on the control plane node that uses the nginx:1.17

```
controlplane $  
controlplane $ cd /etc/kubernetes/manifests/  
controlplane $ pwd  
/etc/kubernetes/manifests  
controlplane $  
controlplane $ ls  
etcd.yaml kube-apiserver.yaml kube-controller-manager.yaml kube-scheduler.yaml  
controlplane $  
controlplane $ k run static-control --image=nginx:1.17 --nodeName=controlplane --dry-run=client  
error: unknown flag: --nodeName  
See 'kubectl run --help' for usage.  
controlplane $ k run static-control --image=nginx:1.17 --dry-run=client  
pod/static-control created (dry run)  
controlplane $  
controlplane $ k run static-control --image=nginx:1.17 --dry-run=client -o yaml > static-control.yaml  
controlplane $  
controlplane $ ls  
etcd.yaml kube-apiserver.yaml kube-controller-manager.yaml kube-scheduler.yaml static-control.yaml  
controlplane $  
controlplane $ nano static-control.yaml  
controlplane $  
controlplane $ k get pods -A -o | grep static  
error: flag needs an argument: 'o' in -o  
See 'kubectl get --help' for usage.  
controlplane $ k get pods -A -o wide | grep static  
default          static-control-controlplane          1/1      Running    0          66s      192.168.0.4    controlplan  
e <none>          <none>  
controlplane $  
controlplane $
```

Q16

Create a new PersistentVolume named safari-pv. It should have a capacity of 2Gi, accessMode ReadWriteOnce, hostPath /Volumes/Data and no storageClassName defined.

Next create a new PersistentVolumeClaim in Namespace project-tiger named safari-pvc . It should request 2Gi storage, accessMode ReadWriteOnce and should not define a storageClassName. The PVC should bound to the PV correctly.

Finally create a new Deployment safari in Namespace project-tiger which mounts that volume at /tmp/safari-data. The Pods of that Deployment should be of image httpd:2.4.41-alpine.

```
controlplane $  
controlplane $ nano q16-pv.yaml  
controlplane $  
controlplane $ k apply -f q16-pv.yaml  
persistentvolume/safari-pv created  
Error from server (NotFound): error when creating "q16-pv.yaml": namespaces "project-tiger" not found  
controlplane $  
controlplane $ k create ns project-tiger  
namespace/project-tiger created  
controlplane $  
controlplane $ k apply -f q16-pv.yaml  
persistentvolume/safari-pv configured  
persistentvolumeclaim/safari-pvc created  
controlplane $  
controlplane $ k get pvc -n project-tiger  
NAME          STATUS    VOLUME    CAPACITY   ACCESS MODES   STORAGECLASS   VOLUMEATTRIBUTESCLASS   AGE  
safari-pvc    Bound     safari-pv  2Gi        RWO            <unset>         <unset>                 8s  
controlplane $
```

```

GNU nano 4.8
apiVersion: v1
kind: PersistentVolume
metadata:
  name: safari-pv
spec:
  storageClassName: ""
  capacity:
    storage: 2Gi
  accessModes:
    - ReadWriteOnce
  hostPath:
    path: "/volumes/data"

---

apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: safari-pvc
  namespace: project-tiger
spec:
  storageClassName: ""
  accessModes:
    - ReadWriteOnce
  resources:
    requests:
      storage: 2Gi

```

```

GNU nano 4.8
apiVersion: apps/v1
kind: Deployment
metadata:
  creationTimestamp: null
  labels:
    app: safari
  name: safari
  namespace: project-tiger
spec:
  replicas: 1
  selector:
    matchLabels:
      app: safari
  strategy: {}
  template:
    metadata:
      creationTimestamp: null
      labels:
        app: safari
    spec:
      volumes:
        - name: deploy-pv
          persistentVolumeClaim:
            claimName: safari-pvc
      containers:
        - image: httpd:2.4.41-alpine
          name: httpd
          volumeMounts:
            - mountPath: "/tmp/safari-data"
              name: deploy-pv
          resources: {}
status: {}

```

```

controlplane $
controlplane $ k create deployment safari --image=httpd:2.4.41-alpine -n project-tiger --dry-run=client
deployment.apps/safari created (dry run)
controlplane $
controlplane $ k create deployment safari --image=httpd:2.4.41-alpine -n project-tiger --dry-run=client > safari.yaml
controlplane $
controlplane $ nano safari.yaml
controlplane $ k create deployment safari --image=httpd:2.4.41-alpine -n project-tiger --dry-run=client -o yaml > safari.yaml
controlplane $
controlplane $ nano safari.yaml
controlplane $
controlplane $ k apply -f safari.yaml
deployment.apps/safari created
controlplane $
controlplane $ k get deployments.apps safari -n project-tiger
NAME      READY   UP-TO-DATE   AVAILABLE   AGE
safari    1/1     1            1           14s
controlplane $

```



```

controlplane $ k describe deployments.apps safari -n project-tiger
Name:          safari
Namespace:     project-tiger
CreationTimestamp: Sun, 22 Sep 2024 03:29:15 +0000
Labels:        app=safari
Annotations:    deployment.kubernetes.io/revision: 1
Selector:      app=safari
Replicas:      1 desired | 1 updated | 1 total | 1 available | 0 unavailable
StrategyType:   RollingUpdate
MinReadySeconds: 0
RollingUpdateStrategy: 25% max unavailable, 25% max surge
Pod Template:
  Labels:  app=safari
  Containers:
    httpd:
      Image:      httpd:2.4.41-alpine
      Port:       <none>
      Host Port:  <none>
      Environment: <none>
      Mounts:
        /tmp/safari-data from deploy-pv (rw)
  Volumes:
    deploy-pv:
      Type:      PersistentVolumeClaim (a reference to a PersistentVolumeClaim in the same namespace)
      ClaimName: safari-pvc
      ReadOnly:  false
  Node-Selectors: <none>
  Tolerations:    <none>
Conditions:
  Type           Status  Reason
  ----           -
  Available      True    MinimumReplicasAvailable
  Progressing    True    NewReplicaSetAvailable
OldReplicaSets: <none>
NewReplicaSet:  safari-668f9f5dc6 (1/1 replicas created)
Events:

```

```

controlplane $
controlplane $ k get pods -n project-tiger
NAME                                READY   STATUS    RESTARTS   AGE
safari-668f9f5dc6-ffr96            1/1     Running   0           86s
controlplane $
controlplane $

```

Q17

Create a NodePort service to expose a pod named **my-pod** on port 8080, with the NodePort set to **30080**.

```
controlplane $  
controlplane $ k run my-pod --image=nginx  
pod/my-pod created  
controlplane $  
controlplane $ k get pods  
NAME                READY   STATUS    RESTARTS   AGE  
my-pod               1/1     Running   0           62s  
static-control-controlplane  1/1     Running   0           13m  
web-server           1/1     Running   0           19m  
controlplane $
```

```
controlplane $  
controlplane $ k create svc nodeport my-pod-nodeport --tcp=8080:8080 --node-port=30080  
service/my-pod-nodeport created  
controlplane $  
controlplane $ k get svc  
NAME                TYPE        CLUSTER-IP   EXTERNAL-IP   PORT(S)          AGE  
kubernetes           ClusterIP   10.96.0.1    <none>         443/TCP          12d  
my-pod-nodeport      NodePort    10.111.163.75 <none>         8080:30080/TCP   3s  
controlplane $
```